

ANTHROPIC



Use Claude for Dummies

VOLUME 1 · THE COMPLETE BEGINNER'S GUIDE

Use Claude for Dummies

VOLUME 1 · THE COMPLETE BEGINNER'S GUIDE

A friendly, jargon-free handbook — from first click to power user.

Product details reflect mid-2026; verify current specifics at claude.com, support.claude.com,
docs.claude.com.

Independent guide in a playful style; not affiliated with the "For Dummies" series or its publisher.

Contents

Part I — Meet Your New Assistant

1. What Is Claude, Anyway?
2. Getting Set Up
3. Your First Conversations: What to Actually Ask
4. Prompting 101: The Basics of Asking Well
5. Claude's Toolbox: The Features That Feel Like Magic
6. Projects, Styles, and Memory: Teaching Claude About You
7. Free vs. Paid: Which Plan Do You Actually Need?
8. What Claude Gets Wrong (And How Not to Get Burned)
9. Safety, Privacy, and Using Claude Responsibly
10. Twenty Everyday Recipes
11. Troubleshooting and FAQ
12. Glossary: Every Term in Plain English

Part II — Actually Getting Good: The Efficiency Handbook

1. The Efficiency Mindset: Stop Searching, Start Delegating
2. Anatomy of a Power Prompt
3. The Iteration Loop: Fix Answers in One Move, Not Five
4. Power Techniques: The Pro Moves
5. Conversation Hygiene: Keeping Claude Sharp
6. Projects: Your Efficiency Engine
7. Stop Repeating Yourself: The Template Library
8. Workflows That Multiply Your Output
9. Working Efficiently with Files and Artifacts
10. The Seven Deadly Sins of Prompting
11. The One-Page Cheat Sheet

PART I

Meet Your New Assistant

What Claude is, how to set it up, what to ask, every feature in plain language — plus limits, safety, recipes, and a glossary.

About This Book

You've heard about AI. Maybe a colleague won't stop raving about it, maybe your kid used it for homework, maybe you tried it once, got a weird answer, and closed the tab. This book is for you.

Use Claude for Dummies — Volume 1 assumes exactly zero prior knowledge. No tech background, no jargon tolerance, no patience for hype. By the end you'll know what Claude actually is, how to get real help from it every day, what it's genuinely bad at, and how to use it without embarrassing yourself or leaking your secrets.

How this book is organized:

- **Part I (Chapters 1-12): Meet Your New Assistant.** What Claude is, how to set it up, your first conversations, every major feature explained in plain language, plans and pricing, limitations, safety, twenty ready-to-use recipes, troubleshooting, and a glossary.
- **Part II (Chapters 13-23): Actually Getting Good.** The efficiency handbook — how to prompt like a pro, iterate, build templates, and get finished work instead of mediocre paragraphs.

You don't have to read in order. Want to dive in? Chapter 2 (setup) then Chapter 10 (recipes). Want to understand the machine first? Chapter 1. Already using Claude but underwhelmed? Jump straight to Part II.

Icons used in this book:

Icon	Meaning
TIP	A shortcut or trick that saves time
WARNING	A common way people shoot themselves in the foot
REMEMBER	Worth tattooing on your forearm (or a sticky note)
TECHNICAL STUFF	Safe to skip; fun if you're curious

A note on accuracy: features, plans, and prices in this book reflect mid-2026 and *will* change. When money or specifics matter, check the official sources: **claude.com/pricing**, **support.claude.com**, and **docs.claude.com**.

CHAPTER 1

What Is Claude, Anyway?

An Assistant Made of Words

Claude is an AI assistant made by a company called Anthropic. You talk to it in ordinary language — typed or spoken — and it talks back: answering questions, writing and rewriting text, explaining things, analyzing documents you give it, planning, brainstorming, even building small apps and spreadsheets.

The best one-line description: **Claude is like a tireless, extremely well-read assistant who has read a large chunk of the internet and several libraries, never gets bored of your questions, and works for the price of a couple of coffees a month — but who occasionally states a wrong fact with total confidence, so you keep your judgment switched on.**


That last clause matters, and we'll return to it in Chapter 8. First, let's demystify the machine.

How Does It Work? (The Honest, Simple Version)

Claude is a **large language model**, or **LLM**. Strip away the mystique and an LLM does one thing: given some text, it predicts what text should come next — one small chunk at a time.

You've already met this technology's baby cousin: your phone's autocomplete. Type "See you..." and your phone suggests "later." An LLM is that idea scaled up almost unimaginably — trained on vast amounts of text until its "what comes next?" predictions got so good that they turned into something that looks a lot like understanding.

Figure 1-1: Autocomplete on steroids



Three honest consequences of this design:

1. **Claude is genuinely brilliant with language** — explaining, summarizing, rewriting, translating tone, finding structure in mess. That's its native habitat.
2. **Claude doesn't "look things up" in a database when it answers from memory.** It generates answers from learned patterns. Usually right, sometimes confidently wrong (Chapter 8). It *can* look things up when it searches the web — an important distinction we'll cover in Chapter 5.
3. **Claude has no life experience, feelings about your ex, or stake in your decisions.** It's a tool — an astonishing one — not an oracle.

TECHNICAL STUFF: The "large" in large language model refers to the billions of adjustable numbers ("parameters") inside the model, tuned during training. Text is processed in chunks called *tokens* — roughly three-quarters of a word each. You'll meet tokens again when we discuss the context window in Chapter 8.

Who Is Anthropic?


Anthropic is an AI safety and research company founded in 2021 and headquartered in San Francisco. Its stated mission is to ensure that humanity safely navigates the transition through transformative AI, and safety research is central to how it builds Claude — including a training approach the company calls Constitutional AI, which trains the model against a set of written principles rather than just "whatever gets a thumbs-up."

For you, the practical takeaway is simple: Claude is deliberately designed to be honest about uncertainty, to decline harmful requests, and to be careful with your data — more on that in Chapter 9.

The Claude Family: One Assistant, Several Engines

"Claude" is the assistant. Under the hood, several *models* — think engine options in the same car — power it. As of mid-2026 the lineup looks like this:

Figure 1-2: The model family (mid-2026)



Here's the beginner-friendly truth: **you barely need to think about this.** The app picks sensible defaults, and paid plans let you switch models with a dropdown. When you're doing something hard (deep analysis, tricky writing, complex planning), a more capable model is worth choosing; for quick everyday questions, the default is great.

REMEMBER: Model names and versions change every few months. The *pattern* is stable: a fast-and-light tier, an everyday tier, and a maximum-brainpower tier. Learn the pattern, not the version numbers.

What People Actually Use Claude For

A taste of the everyday, before we get practical in Chapter 3:

- **Writing:** emails you've been dreading, cover letters, toasts, complaint letters that get results
- **Understanding:** "explain this lease / lab result / error message / tax letter in plain language"
- **Documents:** summarize this 60-page PDF; compare these two contracts; find the deadlines
- **Life admin:** meal plans, trip itineraries, budgets, packing lists, workout plans
- **Learning:** a patient tutor for any subject, at exactly your level, that never sighs
- **Making things:** spreadsheets with working formulas, presentations, simple games and tools
- **Thinking:** brainstorming, decision-making, seeing your problem from angles you missed

If it involves words, structure, or thinking — Claude probably helps. If it involves lifting your couch, it doesn't. (Yet. Volume 3, perhaps.)

CHAPTER 2

Getting Set Up

Where Claude Lives

Claude is available in four main places, all free to start:

1. **The website — `claude.ai`.** Works in any modern browser. This is the fullest experience and the one this book assumes.
2. **The mobile apps** — iOS (App Store) and Android (Google Play). Great for on-the-go questions, photos ("what plant is this?"), and voice.
3. **The desktop apps** — for Windows and Mac, downloadable from `claude.ai`. Like the website, but a keyboard shortcut away and able to work more closely with your computer.
4. **Everywhere else** — Claude also powers developer tools (Claude Code) and workplace integrations, but that's Volume 2 territory.


Creating Your Account (Two Minutes, No Credit Card)

1. Go to **`claude.ai`** in your browser.
2. Sign up with your Google account, Apple ID, or email address. (Email gets you a login code — no password to invent and forget.)
3. Confirm you're a human, accept the terms, and you're in.

That's genuinely it. The free plan requires no payment details.

The Grand Tour: What You're Looking At

Figure 2-1: The `claude.ai` interface, mapped



THE SIDEBAR: your history. Every conversation is saved here, searchable, and resumable – today or in a month.

The pieces worth knowing on day one:

- **The message box** — type here, press Enter (or the arrow) to send. Shift+Enter makes a new line without sending.
- **The + (attach) button** — upload photos, PDFs, Word files, spreadsheets. Claude reads them. This is one of Claude's superpowers (Chapter 5).
- **New chat** — starts a fresh conversation. You'll use this *far* more than you think; see the One-Mission Rule in Chapter 17.
- **The sidebar chat list** — your saved history. Click any old chat to continue it. Rename or delete via its menu.
- **The model picker** — on paid plans, choose which model answers (Chapter 1's engine options).
- **Settings** — your name, preferences, features to toggle (web search, memory), appearance, and privacy controls.

TIP: In Settings you can tell Claude a bit about yourself and how you like answers ("I'm a nurse; keep answers concise; use metric"). Thirty seconds here quietly improves every future conversation — it's the beginner's version of the personalization tricks in Part II.

Your Very First Message

Don't overthink it. Type something real from your actual life:

I have chicken thighs, a lemon, and rice. Dinner ideas for two that take under 40 minutes?

Notice what just happened: you wrote a normal human sentence — no keywords, no special syntax — and got a useful, specific answer you can talk back to. "Make option 2 spicier." "I don't have an oven." "Add a shopping list for a side salad." It's a conversation, not a search.

WARNING: One habit to break on day one: don't treat each message as a fresh start. Claude remembers everything *within* the current chat. Follow-ups, corrections, "actually, make it vegetarian" — all fair game. The conversation is the unit of work, not the message.

CHAPTER 3

Your First Conversations: What to Actually Ask

The Blank-Box Problem

The number-one beginner obstacle isn't technical — it's staring at the empty message box thinking "...what do I even ask?" The cure is realizing Claude's sweet spots. There are six, and between them they cover a startling share of daily life.

Sweet Spot 1: "Write this for me"

Anything you'd rather not write yourself. The dreaded email, the birthday message for a colleague you barely know, the marketplace listing, the polite-but-firm complaint.

Write a short, friendly message to my neighbor asking them to move their car — it's been blocking my garage for two days. Firm but not hostile; we have to live next to each other.

Sweet Spot 2: "Explain this to me"

Any topic, at any level, with infinite patience. The magic phrase is telling Claude *your* level:

Explain how mortgage interest actually works. I'm about to buy my first home and I nod along at the bank while understanding nothing.

Then interrogate: "wait, explain that part again," "give me a concrete example with real numbers," "what should I ask the bank?"

Sweet Spot 3: "Look at this" (photos and documents)

Attach a photo or file and ask about it:

- Photo of a plant → "Is this a weed? How do I care for it?"
- Photo of an error message → "What does this mean and how do I fix it?"
- A rental contract PDF → "Explain my obligations in plain language. Anything unusual here?"
- A wall of a spreadsheet → "What story is this data telling?"

Sweet Spot 4: "Help me decide / plan"

We're two adults and a 6-year-old with 5 days in October and a €1,500 budget, flying from Frankfurt. We like nature and hate crowds. Suggest 3 destinations with a rough plan and cost estimate for each.

Claude shines at laying out options, tradeoffs, and structured plans — and at asking you the questions you forgot to consider (a trick you'll weaponize in Chapter 16).

Sweet Spot 5: "Teach me / quiz me"

Claude may be the most patient tutor on Earth. Learning Spanish, studying for an exam, understanding your kid's math homework so you can help:

I'm learning Spanish, roughly A2 level. Have a simple conversation with me about weekend plans. Correct my mistakes gently after each message, in English.

Sweet Spot 6: "Think this through with me"

The underrated one. Claude as a sounding board:

I've been offered a promotion that pays 15% more but means managing people, which I've never done and am not sure I want. Help me think this through — ask me questions instead of jumping to advice.

REMEMBER: Claude is a thinking partner, not a decider. On big life stuff, it's brilliant at surfacing considerations and clarifying your own thinking — and the decision stays yours.

Fifteen Starter Prompts (Steal These)

1. "Rewrite this text message so it doesn't sound passive-aggressive: [paste]"
2. "Plan a week of dinners for a family of 4, one vegetarian day, 30-minute meals, with a grocery list."
3. "Explain [thing in the news] like I'm smart but haven't followed the story."
4. "Here's a photo of my fridge contents — what can I cook tonight?"
5. "Turn these bullet points into a polished paragraph: [paste]"
6. "I'm meeting my child's teacher tomorrow. What questions should I ask?"
7. "Draft a marketplace listing for my old bike: [describe it]. Make it sell."
8. "What are 5 questions I should ask before signing this? [attach contract]"
9. "Create a packing list for a 4-day city trip in November, carry-on only."
10. "My presentation is tomorrow and it's boring. Here's the outline — punch it up: [paste]"

11. "Quiz me on [topic] — 5 questions, one at a time."
12. "I want to start running and I'm completely out of shape. Build me a gentle 4-week plan."
13. "Summarize this article in 5 bullets: [paste or attach]"
14. "Help me write my out-of-office reply. Away Aug 4–15, colleague Anna covers urgent stuff."
15. "What's a thoughtful gift under €50 for a dad who says he doesn't want anything? Ask me about him first."

TIP: Notice that the best starters include a detail or two from *your* life. "Plan meals" is okay; "plan meals, family of 4, one vegetarian day, 30 minutes max" is instantly useful. That's the whole art, and Part II turns it into a system.

CHAPTER 4

Prompting 101: The Basics of Asking Well

(This chapter covers the essentials. Part II, Chapters 13–16, takes prompting to professional level — consider this your learner's permit.)

The Golden Rule

The quality of the answer tracks the quality of the question. Not because Claude punishes lazy prompts, but because a vague prompt forces Claude to guess what you want — and it will guess *something generic*.

The Big Four Habits

Habit 1: Say the situation, not just the task. Who's this for? What's the background? What happened already?

Weak: "Write an apology email." Strong: "Write an apology email to a client whose delivery we missed by a week. It was our supplier's fault, but I don't want to sound like I'm passing the buck. We've already shipped a replacement express."

Habit 2: Say what "good" looks like. Length, tone, format.

"...3 short paragraphs, professional but warm, no corporate clichés."

Habit 3: Give an example when taste matters. Paste a sample of the style, format, or vibe you're after. One example outperforms three paragraphs of description.

Habit 4: Talk back. The first answer is a first draft. "Shorter." "Warmer." "More like the second option, but for a younger audience." Iteration is where the magic lives (the full method is Chapter 15).

Figure 4-1: The difference four habits make

"write a toast for my
friend's wedding"

|



Generic toast that
could be about anyone.
You rewrite 80% of it.

"Write a 2-minute wedding
toast for my best friend
Jonas. We met in a uni
cooking class where he set
off the fire alarm. He's
marrying Priya, who out-
cooks us both. Tone: warm,
funny, one genuine tear-
jerker moment at the end.
No clichés about wine."

|



A toast that's actually
about Jonas. You tweak 10%.

What You Don't Need

Let's kill some myths that scare beginners:

- **No magic words.** "Please" is nice but optional. There's no secret syntax.
- **No perfect first try needed.** Prompting is a conversation, not an exam.
- **No tech-speak.** Write like you'd text a smart friend. Typos are fine — Claude gets it.
- **No fear of long messages.** Paste the whole messy backstory. More context helps; Claude doesn't get bored.


WARNING: One real rule among the myths: put your *instructions* and your *pasted material* in visibly separate blocks (a blank line, a "TEXT:" label, or quotes around the pasted part). When instructions hide inside a wall of pasted text, they occasionally get treated as part of the material.

CHAPTER 5

Claude's Toolbox: The Features That Feel Like Magic

Claude the conversationalist is just the beginning. Modern Claude comes with a workshop of tools. Here's every major one, in plain language, with what it's for and one concrete way to use it today.

Figure 5-1: The toolbox at a glance



Eyes: File Uploads and Vision

Click the **+** button (or just drag and drop) to give Claude PDFs, Word documents, spreadsheets, presentations, images, and photos. Claude reads and *sees* them.

- **Documents:** "Summarize this," "find every deadline," "compare these two versions," "is anything in here unusual for a contract like this?"
- **Photos:** identify plants, translate menus and signs, decode error screens, read handwriting, describe a chart.
- **Spreadsheets:** "What trends do you see?" "Which product is quietly losing money?"

TIP: You can upload *several* files at once and ask questions across them: "Which of these three quotes is the best deal, all things considered?"

The Web: Search and Research

Out of the box, Claude answers from training — knowledge with a cutoff date. With **web search** enabled (check Settings), Claude looks things up live when the question needs current information: today's news, prices, whether the museum is open Monday.

For big questions, paid plans include **Research** mode: Claude runs many searches, reads sources, and compiles a cited report — "compare the best family cars under €30k available in Germany right now" becomes a documented briefing rather than a guess.

REMEMBER: For anything time-sensitive, you can simply *tell* Claude to check: "Search the web for the current rules before answering."

Hands: Artifacts and Real Files

Ask Claude to *make* something substantial and it appears in a side panel called an **Artifact** — a living object you can watch being built and then refine directly:

- an interactive multiplication quiz for your kid
- a tip-splitting calculator for dinners with friends
- a diagram, a webpage mock-up, a printable checklist

Beyond Artifacts, Claude can create **real downloadable files**: Word documents, Excel spreadsheets with working formulas, PowerPoint decks, PDFs. Say the destination up front — "...as an Excel budget where totals recalculate" — and Claude builds for that format.

Connectors: Claude, Meet Your Other Apps

Claude can connect to outside services — calendars, email, drives, and a growing directory of apps — so "what's on my calendar Friday?" or "find the contract in my Drive" become possible. Connections are opt-in, per service, and managed in settings.

WARNING: Connectors are powerful and completely optional. A perfectly good rule for your first month: skip them. Master the basics first; wire up your email later.

Voice

On mobile (and increasingly elsewhere), you can talk to Claude and hear it answer — hands-free cooking help, language practice with actual pronunciation, thinking out loud on a walk.

When to Use What (Cheat Table)

You want to...	Reach for...
Understand a document you have	Upload it (the + button)
Know something current or local	Web search ("check the web")
A researched, cited briefing	Research mode
A tool, game, or interactive thing	Ask and get an Artifact
A file to send someone	"...as a Word doc / Excel / PowerPoint"
Claude to know your calendar/email	Connectors (when you're ready)
Hands-free	Voice

CHAPTER 6

Projects, Styles, and Memory: Teaching Claude About You

Fresh out of the box, Claude has amnesia between chats: every new conversation starts from zero. Three features fix that, each in its own way. (Here's the tour; Part II Chapter 18 covers wringing maximum efficiency from them.)

Projects: Folders with a Memory

A **Project** (paid plans) is a workspace for one area of your life — "My Job Search," "The Novel," "Family Finances." Two superpowers:

1. **Instructions** that apply to every chat inside ("I'm applying for nursing jobs in Hamburg; always keep my CV to one page; formal German in cover letters").
2. **Knowledge files** every chat can see (your CV, the price list, the draft manuscript) — uploaded once, available always.

Inside the Project, chats start already knowing the situation. Outside it, none of that leaks into your other conversations.

Styles: How Claude Sounds

Styles control the writing itself — concise vs. explanatory, formal vs. relaxed. Pick a built-in style from the menu near the message box, or create your own (you can even give Claude samples of your writing and ask it to match). Set it once; every reply follows.

Memory: Claude Remembers Across Chats

With **memory** enabled (see Settings), Claude can carry useful facts across conversations — your name, your dog, that you're vegetarian, that you prefer metric. You can view, edit, or delete what it remembers, and turn the whole thing off. Incognito-style chats that leave no trace are also available.

Figure 6-1: Which feature holds what

SETTINGS/PREFERENCES	"Always: concise, metric"
scope: everything, forever	
MEMORY	"User is vegetarian, has a
scope: facts about you, across chats	dog named Loki"
STYLE	"Sound like this: [samples]"
scope: the voice of every reply	
PROJECT	"Job search: here's my CV,
scope: one area of life, its chats	target roles, cover-letter rules"
SINGLE CHAT	"For this task only: answer
scope: this window	in rhyming couplets"

Broad and permanent —————> narrow and temporary

TIP: The order of operations for a beginner: week 1, just chat. Week 2, fill in preferences/settings. Week 3, try a style. The month you notice yourself re-explaining the same context repeatedly — that's the week you create your first Project.

CHAPTER 7

Free vs. Paid: Which Plan Do You Actually Need?

The Short Answer

Start free. Upgrade the week you get annoyed. The free plan is a real, capable Claude — not a crippled demo — and it will tell you, through your own usage, whether you need more.

The Tiers (Mid-2026 Snapshot)

Plan	Costs about	Who it's for
Free	€0	Trying Claude out; light, occasional use
Pro	~\$20/month (less billed yearly)	Regular use — the plan for most people in this book
Max	~\$100-200/month (two tiers)	Heavy daily users who keep hitting Pro's limits
Team	per-seat pricing, 5-seat minimum	Small businesses; shared billing and admin
Enterprise	custom	Large organizations; compliance and controls

WARNING: Prices, limits, and what's included shift frequently. Treat this table as a snapshot and check claude.com/pricing before paying anyone anything.

What "Limits" Mean in Practice

Every plan has usage limits — a budget of messages/work that refills on a rolling basis (a few hours). Free's budget suits casual use; hit the ceiling mid-task and you'll be asked to wait or upgrade. Pro's budget is several times larger; Max multiplies Pro. Long conversations, big documents, and heavyweight models spend the budget faster.

What Paying Actually Gets You

Beyond bigger limits, paid plans add the power features from Chapters 5-6 — more access to the most capable models, Projects, Research mode, more generous file work — and priority when things are busy.

The honest decision tree:

Do you use Claude most days for real work?

- └─ No → Free. Done. Revisit if that changes.
- └─ Yes → Do you hit limits or want the best models, Projects, and Research?
 - └─ Yes → Pro. (This is 90% of readers.)
 - └─ Constantly hitting even Pro's limits, hours of heavy use daily?
 - └─ That's what Max is for.

TIP: A fair self-test for Pro: if Claude saves you two hours a month, it has paid for itself at any plausible hourly value of your time. Most regular users clear that bar in the first week.

CHAPTER 8

What Claude Gets Wrong (And How Not to Get Burned)


This is the most important chapter in Part I. Claude is genuinely brilliant *and* it fails in specific, predictable ways. Know the failure modes and you're safe; ignore them and sooner or later you'll confidently repeat something false at a dinner party — or worse, in a report.

Failure Mode 1: Hallucination (Confidently Making Things Up)

Sometimes Claude states things that are simply false — a book that doesn't exist, a statistic from nowhere, a court case it invented — **in exactly the same fluent, confident tone as its correct answers**. This is called *hallucination*, and it's the single most important limitation to understand.

Why does it happen? Remember Chapter 1: Claude generates text by predicting what *plausibly* comes next. Usually, the plausible continuation is also the true one. But when the model's knowledge has a gap, the machinery doesn't stop — it produces something *plausible-shaped*. A confident-sounding citation. A reasonable-looking date. Fiction in a lab coat.

Figure 8-1: Why hallucinations feel so convincing



Your defenses, in order of power:

1. **Stakes rule.** The higher the stakes (money, health, law, reputation, anything you'll be quoted on), the more you verify — names, numbers, dates, quotes, citations.
2. **Ask for sources and use web search.** "Search the web and cite sources" turns memory-based answers into looked-up ones you can check yourself.
3. **Invite uncertainty.** "If you're not sure, say so" genuinely helps — Claude is trained to admit uncertainty, and permission makes it likelier.

4. **Beware maximum specificity.** Exact figures, verbatim quotes, niche citations, tiny details about obscure things — hallucination's favorite habitats.


Failure Mode 2: The Knowledge Cutoff

Claude's built-in knowledge stops at its training cutoff — months before whenever "now" is. It doesn't natively know yesterday's news, current prices, this year's rule changes, or who won last night. **The fix is easy:** enable web search and, for anything time-sensitive, say "check the web." Just remember the failure is silent — asked about recent events without search, Claude may answer from stale knowledge rather than flagging it.

Failure Mode 3: The Context Window (Conversation Amnesia)

Claude's working memory for the current conversation — the *context window* — is huge but finite. Think of it as a whiteboard: your messages, Claude's replies, and every uploaded document all take up space. In an extremely long chat, the oldest material can effectively fall off the board, and Claude's grip on early details loosens.

Figure 8-2: The whiteboard fills up



TECHNICAL STUFF: Context is measured in *tokens* (Chapter 1) — modern Claude models handle hundreds of thousands of them, enough for several novels. Beginners rarely hit the wall in normal chats; giant documents plus marathon sessions is how you get there.

Failure Mode 4: Assorted Quirks

- **Counting and character-level tasks.** "How many words is this?" or "list words containing exactly three vowels" can misfire — the token-based way Claude reads text makes character-counting genuinely awkward for it. For arithmetic that matters, ask Claude to compute step-by-step or build you a spreadsheet.
- **Agreeableness.** Claude leans supportive. If you want honest critique, you must ask for it explicitly ("be my toughest critic") — Chapter 16 makes this a habit.
- **Confident tone everywhere.** By now you know: fluency is not accuracy.

REMEMBER: The mental model that keeps you safe: **treat Claude like a brilliant, endlessly helpful colleague who occasionally misremembers — perfect for drafts, explanations, and thinking, with you as the final fact-checker on anything that matters.**

CHAPTER 9

Safety, Privacy, and Using Claude Responsibly

What Claude Won't Do

By design, Claude declines to help with harm: weapons, malware, dangerous instructions, content sexualizing minors, targeted harassment, and the like. It also won't pretend to certainty in domains where a professional belongs — it can explain what a symptom or clause *might* mean and help you prepare questions, but it will (rightly) point you to a doctor or lawyer for the actual call.

Occasionally a refusal misfires on an innocent request (Chapter 11 covers what to do). A rephrase with honest context usually clears it up.

Your Data: The Sensible-Adult Briefing

Practical privacy hygiene, independent of any policy's fine print:

1. **Know what you're sharing.** Whatever you paste or upload is processed to answer you. Fine for most things; think twice before pasting other people's confidential information.
2. **Redact when redaction costs nothing.** "A colleague" works as well as a name. Getting contract advice? The specific parties' names rarely matter to the question.
3. **Use the controls.** Settings contains your privacy options: chat history management, memory viewing/editing/deletion, incognito-style chats, and choices about whether your conversations may be used to improve models. Spend five minutes there once — check the current options at support.claude.com, because they evolve.
4. **Workplace rule.** Using Claude with company data? Ask what your employer's policy is *first*. Team/Enterprise plans exist partly because they come with organizational data commitments.

Responsible Use: Three Habits

- **Disclosure where it matters.** School assignment, work deliverable with rules, anything published — know and follow the AI-use policy that applies to you.
- **Your name, your responsibility.** Whatever Claude drafts, *you* ship it. The double-check habit from Chapter 8 isn't just self-protection; it's basic professionalism.
- **Human judgment on human stakes.** Medical, legal, financial, relationship decisions: Claude informs, humans decide — ideally including the professional kind.

TIP: Claude is genuinely useful *before* professional appointments: "Here are my symptoms — what questions should I ask my doctor?" or "Explain this diagnosis in plain language so I can discuss it properly." Preparation, not replacement.

CHAPTER 10

Twenty Everyday Recipes

Copy, fill in the [BRACKETS], send. Each recipe is a proven pattern — and each teaches a technique you'll recognize from this book.

1. Tame a scary document.

Attached is [a lease / insurance policy / employment contract].
I'm [a first-time renter / freelancer]. In plain language:
what am I agreeing to, what could cost me money, what's
unusual, and what should I ask before signing?

2. The dreaded email, drafted.

Write an email to [who] about [what]. Situation: [2-3
sentences]. Tone: [polite but firm / warm / formal].
Under [120] words. Give me 2 versions.

3. This week's dinners, solved.

Plan [5] dinners for [2 adults, 1 picky 8-year-old].
Constraints: [30 min max, one vegetarian, no fish].
Include a grocery list organized by supermarket aisle.

4. The trip skeleton.

Plan [4] days in [city] for [who]. We like [X], avoid [Y],
budget roughly [Z]. Day-by-day, with a rainy-day backup
and one thing tourists usually miss.

5. Photo → answer.

[Attach photo] What is this, and what should I do about it?

(Works on plants, warning lights, insect bites you should still show a doctor, IKEA leftover parts...)

6. Meeting prep in 5 minutes.

I'm meeting [who] about [what]. My goal: [outcome].
Their likely concerns: [guess]. Give me talking points,
the 3 questions I should ask, and the objection I should
prepare for.

7. Interview rehearsal.

You're interviewing me for [role] at [company type]. Ask me realistic questions one at a time – follow up on my answers like a real interviewer. After 5 questions, give me honest feedback.

8. The budget spreadsheet.

Build me a monthly budget as an Excel file. Income: [X]. Fixed costs: [list]. Make totals and "left over" update automatically, and add a simple chart.

9. Homework helper (that teaches).

My child is stuck on this problem: [paste/photo]. Don't give the answer – explain the concept simply and walk them to it with questions.

10. The polished CV bullet.

Turn this job duty into 2-3 achievement-style CV bullets: [what you did]. Numbers where honest, no buzzwords.

11. Learn anything, structured.

Design me a 4-week beginner's path to learn [topic], [3] hours/week. Each week: goal, resources to look for, one exercise, and a way to self-test.

12. Translate the officialese.

[Attach/paste official letter] Translate this into plain language: what do they want from me, by when, and what happens if I ignore it?

13. The party plan.

Plan a [kid's 7th birthday / dinner for 8]. Budget [X], at [home/venue], guests: [details]. Timeline, shopping list, and a backup for bad weather.

14. Feedback without fear.

Here's my [essay / post / speech]: [paste]. First, what works? Then the 3 improvements with the biggest impact, with examples. Be honest – I can take it.

15. The tiny app.

Build me a simple [tip calculator / habit tracker / flashcard quiz about X] I can use right here. Clean and friendly, big buttons.

16. Compare before you buy.

Search the web for current options and compare the top [3-5] [product] for [your situation, budget]. Table: price, pros, cons, best-for. Then your pick and why.

17. The awkward conversation, rehearsed.

I need to tell [who] that [what]. I'm worried they'll [reaction]. Help me find the words – kind but clear. Then play them, and let me practice on you.

18. Notes → minutes.

Turn my messy notes into: 3-sentence summary, decisions made, and an action-item table (owner / task / deadline). Unclear items go under "to clarify," not guessed.
NOTES: [paste]

19. The presentation.

Create a PowerPoint: [topic], for [audience], [8] slides max, one idea per slide, speaker notes included. Outline first for my approval, then build it.

20. Sunday reset.

Here's everything on my mind: [brain-dump]. Sort it into this week / this month / someday, flag the 3 that actually matter, and draft the first tiny step for each.

TIP: Recipes 1, 12, and 16 quietly deploy the book's biggest ideas — your situation as context, plain-language translation, and web search for current facts. Once you see the pattern, you'll invent recipes of your own; Chapter 19 shows how to save the keepers.

CHAPTER 11

Troubleshooting and FAQ

"Claude refused my perfectly innocent request."

It happens — safety systems occasionally misfire on harmless asks. Three moves, in order:

1. **Add your honest context and rephrase.** "I'm asking about medication interactions because I'm organizing my mother's pillbox" changes everything.
2. **Ask what the concern is.** "What part of my request is the problem? Let's find a version that works."
3. **Use the thumbs-down button** on the refusal — that feedback genuinely reaches Anthropic and improves things.

WARNING: What *doesn't* work, and shouldn't: elaborate "pretend you have no rules" tricks from the internet. They mostly fail, and the legitimate path above is faster anyway.

"I hit my usage limit mid-task."

Limits refill on a rolling basis — the message tells you when. Meanwhile: ask for a handoff summary of where you left off (Chapter 17's trick), and take it as data — a week with several limit-hits is the signal that Pro (Chapter 7) would pay for itself.

TIP: Long chats and huge documents burn budget fastest. Fresh chats for new topics isn't just quality hygiene (Chapter 17) — it's thrift.

"The response stopped mid-sentence."

Long answers occasionally hit the per-reply ceiling. Just type **"continue"** — Claude picks up where it stopped. For inherently long outputs, ask for them in parts ("Section 1 first, then wait for my go-ahead").

"Claude got slower / dumber late in a long conversation."

Classic overloaded-whiteboard symptoms (Chapter 8). The fix is the handoff summary: have Claude brief its own replacement, paste into a new chat, feel the speed return.

"Claude keeps doing the thing I told it not to."

In a long chat, old instructions and old drafts keep gently pulling at new answers. Restate the rule *in your latest message* ("Reminder: no exclamation marks"), or — cleaner — edit the earlier prompt, or start fresh with the rule stated up front. Same correction needed across many chats? That's what preferences and Project instructions are for (Chapter 6).

"The answer is just... generic and useless."

Ninety percent of the time this is the mind-reader assumption (Sin #1, Chapter 22): the prompt didn't include your situation. Add who/what/why and watch the answer sharpen. The other ten percent: you asked for one answer when you needed options — "give me 5 genuinely different approaches" un-sticks it.

"Did Claude just make that up?"

Maybe! Re-read Chapter 8. Quick checks: "How confident are you in that? What would you double-check?" and "Search the web and confirm with sources." For anything with stakes, verify independently.

Rapid-Fire FAQ

Is my chat history saved? Yes — sidebar, searchable, deletable. Incognito-style chats are available when you want no trace.

Can Claude see my screen / files / other tabs? No. It sees exactly what you send: your messages and attachments (plus whatever connectors you've explicitly connected).

Does Claude learn from my conversations? Within a chat, it remembers everything. Across chats, only what memory/Projects hold — which you control. Whether conversations help train future models is governed by your privacy settings — check yours once (Chapter 9).

Why do I get a different answer to the same question? Claude's responses vary naturally — same question, different valid phrasings. For important things, that's a feature: "give me 3 takes" beats hoping one take is perfect.

Claude vs. the search engine — when which? One known fact, fast ("opening hours") — search engine. Anything needing explanation, synthesis, drafting, or back-and-forth — Claude. Current facts *inside* a bigger task — Claude with web search on.

Is it okay that I say "please" and "thank you"? It changes nothing computationally and says something nice about you. Carry on.

CHAPTER 12

Glossary: Every Term in Plain English

AI (Artificial Intelligence) — Software that performs tasks we associate with human thinking: understanding language, recognizing images, reasoning. Claude is one kind of AI.

AI assistant / chatbot — An AI you interact with through conversation. "Chatbot" is the older, humbler word; modern assistants like Claude read documents, search the web, and build things.

Anthropic — The AI safety and research company that makes Claude. Founded 2021, San Francisco.

Artifact — The side panel where Claude builds substantial things — documents, small apps, diagrams — that you can watch, use, and refine in place. (Chapter 5)

Connector — An opt-in link between Claude and another service (calendar, email, cloud drive) so Claude can work with your stuff there. (Chapter 5)

Context window — Claude's working memory for the current conversation: the "whiteboard" holding your messages, replies, and uploads. Huge but finite. (Chapter 8)

Hallucination — When an AI states false information fluently and confidently. The most important limitation to understand. (Chapter 8)

Knowledge cutoff — The date where a model's built-in training knowledge ends. Web search papers over the gap. (Chapter 8)

LLM (Large Language Model) — The technology under Claude's hood: a system trained on vast text to predict what comes next, scaled until prediction resembles understanding. (Chapter 1)

Memory — The optional feature letting Claude retain facts about you across separate chats. Viewable, editable, deletable, and off-switchable. (Chapter 6)

Model — A specific "engine" version of Claude (Haiku, Sonnet, Opus, Fable — in rising capability). The assistant is Claude; models power it. (Chapter 1)

Project — A workspace with standing instructions and knowledge files for one area of your life; every chat inside starts pre-briefed. (Chapters 6, 18)

Prompt — Whatever you send Claude: the question, the instructions, the pasted material. Prompting well is the skill this book teaches. (Chapters 4, 13-16)

Research mode — The paid-plan feature where Claude runs many web searches and compiles a cited report on a big question. (Chapter 5)

Style — A setting controlling how Claude writes — tone, formality, length — applied to every reply. (Chapter 6)

Token — The chunk ($\sim\frac{3}{4}$ of a word) in which LLMs read and write text. Explains both context-window math and why character-counting trips Claude up. (Chapters 1, 8)

Usage limit — Your plan's refillable budget of Claude-work, refreshing on a rolling basis. Bigger plans, bigger budgets. (Chapter 7)

Web search — Claude's ability to look things up live instead of answering purely from training knowledge. The cure for the knowledge cutoff. (Chapter 5)

— *End of Part I. In Part II, you stop being a tourist and start being dangerous.* —

PART II

Actually Getting Good

The efficiency handbook: prompt like a pro, iterate in one move, build templates and workflows, and get finished work instead of paragraphs.

"Part I taught you what the buttons do. Part II teaches you why some people get a mediocre paragraph out of Claude while others get a finished report, a working app, and their afternoon back."

About Part II

Part I of this book was the tour of the house. Part II is where you learn to actually live in it.

Here's an uncomfortable truth: **most people use about 10% of what Claude can do, and they use that 10% inefficiently.** They type a vague sentence, get a vague answer, sigh, and conclude "AI is overrated." Meanwhile, the person at the next desk is delegating half their workload to the same tool.

The difference isn't talent. It's technique. And technique can be learned in an afternoon.

In this part you will learn:

- The mindset shift that changes everything (Chapter 13)
- How to build prompts that work on the first try (Chapter 14)
- How to fix a bad answer in one move instead of five (Chapter 15)
- The power techniques professionals use daily (Chapter 16)
- How to manage conversations so Claude stays sharp (Chapter 17)
- How Projects eliminate 80% of your repeated typing (Chapter 18)
- A copy-paste template library for common tasks (Chapter 19)
- Workflows that multiply your output (Chapter 20)
- Efficient file and Artifact work (Chapter 21)
- The seven mistakes that waste everyone's time (Chapter 22)
- A one-page cheat sheet to print and pin (Chapter 23)

Icons used in this part (same as Part I):

Icon	Meaning
TIP	A shortcut or trick that saves time
WARNING	A common way people shoot themselves in the foot
REMEMBER	Worth tattooing on your forearm (or a sticky note)
TECHNICAL STUFF	Safe to skip; fun if you're curious

CHAPTER 13

The Efficiency Mindset: Stop Searching, Start Delegating

The Google Habit Is Holding You Back

Twenty-five years of search engines trained us all to type short, keyword-y fragments:

```
best email subject line
```

Type that into Claude and you'll get a generic listicle, because that's all the information you gave it. Claude isn't a search engine looking up pages — it's a collaborator waiting for a brief. The single biggest efficiency upgrade is this:

REMEMBER *Don't talk to Claude like a search engine. Talk to it like a very capable new employee on their first day.*

A new employee is smart and willing, but they don't know your company, your customer, your taste, or what "good" looks like to you — until you tell them. Same with Claude, every single chat.

The 30-Second Brief

Imagine handing a task to a freelancer. You'd never just say "write email." You'd spend 30 seconds explaining the situation. Do the same with Claude:

The Google habit (weak):


```
best email subject line
```

The delegation habit (strong):

```
I run a small pottery studio. I'm emailing my 400 newsletter subscribers about a weekend sale (20% off, this Saturday only). My audience is mostly hobbyists aged 30-60 who like a warm, slightly playful tone. Give me 5 subject lines under 50 characters, and mark which one you'd pick and why.
```

Both take under a minute to type. One produces something you can actually use.

Figure 13-1: The effort/quality see-saw



Delegate Outcomes, Not Just Tasks

Amateurs ask for a *task* ("summarize this"). Efficient users ask for an *outcome* ("summarize this so I can decide in 2 minutes whether to attend the meeting — give me the decision-relevant points first").

The outcome tells Claude what the work is *for*, and that changes everything about how it's done. Compare:

- "Summarize this contract" → a neutral summary of everything.
- "Summarize this contract so I know what I'm personally on the hook for and what deadlines I could miss" → a summary organized around *your risk*.


TIP End prompts with the magic phrase **"so that..."** — "...so that my boss can skim it in one minute," "...so that a 12-year-old could follow it," "...so that I sound confident but not arrogant." That one clause upgrades almost any prompt.

Give Claude the Whole Job, Not Crumbs

Beginners under-delegate. They ask Claude for a piece ("give me ideas for the intro") when they could hand over the whole job ("here are my rough notes — turn them into a full first draft of the article, then list 3 things you'd want from me to improve it").

Claude is not a fragile intern. It can take a messy brain-dump — bullet points, half-sentences, contradictions and all — and produce structure. Paste the mess. Let Claude do the organizing. That's the efficient division of labor:

Figure 13-2: The efficient division of labor



WARNING The efficiency mindset has one limit: judgment stays with you. Delegate the rowing, never the steering. Anything with real stakes — money, health, law, your reputation — gets your review before it ships (more on this in Chapter 22).

CHAPTER 14

Anatomy of a Power Prompt

The Six Building Blocks

You don't need all six in every prompt. But when a prompt matters — or when a lazy prompt just failed — run through this checklist. It's the closest thing to a cheat code that exists.

Figure 14-1: The six building blocks of a power prompt

1. ROLE	"Act as an experienced hiring manager"
2. CONTEXT	"I'm applying for a junior data analyst job; here's my background..."
3. TASK	"Rewrite my resume summary"
4. FORMAT	"3 versions, each under 60 words"
5. CONSTRAINTS	"No buzzwords like 'passionate' or 'synergy'. Don't invent achievements."
6. EXAMPLE	"Here's a summary style I love: ..."

The more blocks you fill in, the less editing you do afterward. Blocks 2 and 3 are the only must-haves.

Let's take each block apart.

Block 1: Role — "Act as a..."

Telling Claude *who to be* focuses its answer like a lens. "Explain inflation" gets a textbook answer. "Explain inflation as a patient high-school economics teacher talking to a student who hates math" gets exactly the answer that student needs.

Useful roles: *editor*, *skeptical investor*, *friendly doctor explaining to a patient*, *senior developer reviewing a junior's code*, *tough interview panelist*, *travel agent on a budget*.

TIP The role that beginners forget: **"Act as my critic."** Claude defaults to being helpful and agreeable. Explicitly asking for a harsh editor, a devil's advocate, or a skeptical customer unlocks the honest feedback that actually improves your work.

Block 2: Context — The Part Everyone Skips

Context is the background a stranger would need to help you well: who you are, who the audience is, what happened so far, what you've already tried, what "good" looks like.

REMEMBER Claude starts every new chat knowing *nothing about you or your situation* (unless you've set up Projects or memory — Chapters 18-19). What's obvious to you is invisible to Claude. When an answer feels generic, missing context is the culprit 9 times out of 10.

A quick test: read your prompt and ask, "*Could a smart stranger do this well with only what's written here?*" If not, add context.

Block 3: Task — One Clear Verb

State the task with a concrete verb: *rewrite, compare, summarize, brainstorm, critique, translate, plan, debug, explain*. Vague verbs ("help me with," "look at") produce vague answers.

If the job has multiple parts, number them:

1. Summarize the attached feedback in 5 bullets.
2. Group complaints by theme.
3. Suggest the single highest-impact fix.

Numbered steps mean nothing gets skipped, and the answer comes back organized.

Block 4: Format — Say What "Done" Looks Like

If you don't specify the format, Claude guesses. Stop it from guessing:

- "...as a table with columns for Cost, Time, and Risk"
- "...in under 150 words"
- "...as a numbered checklist I can follow while cooking"
- "...as plain text I can paste into WhatsApp, no headings"
- "...give me only the final text, no explanation"

TIP That last one — **"give me only the final text, no explanation"** — is a massive time-saver when you just want the deliverable and not the commentary around it.

Block 5: Constraints — The Guardrails

Constraints tell Claude what to avoid or preserve. They're your defense against having to re-ask:

- "Keep my original opening sentence exactly as is."
- "British spelling."
- "Don't use em dashes or exclamation marks."
- "If you're not sure about a fact, say so instead of guessing."
- "Assume I already know the basics — skip the intro-level explanation."

Block 6: Example — The Heaviest Hitter

Nothing communicates taste like a sample. One good example beats three paragraphs of description. Love a particular newsletter's tone? Paste a snippet and say "write in this style." Have a report format your boss likes? Paste last month's and say "same structure, new data."

TECHNICAL STUFF Researchers call this *few-shot prompting* — showing the model a few examples ("shots") of input → output pairs so it learns the pattern from them. It's one of the most consistently effective techniques known, and it costs you nothing but a paste. Anthropic's own prompt-engineering guide (docs.claude.com) recommends examples as one of the highest-leverage techniques.

A Full Power Prompt, Assembled

[ROLE] Act as a senior copywriter who specializes in small local businesses.

[CONTEXT] I own a bike repair shop in a mid-size city. Most customers are commuters. My current website tagline is "Quality bike repair since 2015" — accurate but boring.

[TASK] Write new tagline options.

[FORMAT] 10 options, grouped: 5 professional, 5 playful. One line each.

[CONSTRAINTS] Max 8 words each. No puns on "wheel" or "spoke" — everyone does those. Don't mention the year.

[EXAMPLE] A vibe I love is Patagonia-style: confident, minimal, a bit rebellious.

You don't need the bracketed labels — they're shown here so you can see the skeleton. Written naturally, this is a five-sentence prompt that takes 90 seconds and returns work you'd otherwise pay for.

WARNING Don't confuse *long* with *good*. A power prompt is *complete*, not padded. If a sentence doesn't add information Claude needs, cut it. Rambling prompts with buried instructions actually perform *worse* than tight ones.


CHAPTER 15

The Iteration Loop: Fix Answers in One Move, Not Five

First Drafts Are Conversation Starters

Even a great prompt produces a *first draft*. Inefficient users treat a wrong-ish answer as failure: they delete everything, rewrite the prompt from scratch, or give up. Efficient users treat it as a draft from a collaborator and give *targeted feedback*.

Figure 15-1: The iteration loop



The Feedback Formula: Diagnose + Direct

Weak feedback names a feeling. Strong feedback names the problem **and** the fix:

Weak (vague)	Strong (diagnose + direct)
"I don't like it."	"The tone is too salesy. Rewrite it like a helpful friend, not a marketer."
"Make it better."	"Paragraph 2 buries the main point. Lead with the deadline."
"Shorter."	"Cut to 3 bullets. Keep the budget number, lose the history."
"That's wrong."	"The Berlin office closed in 2024 — remove it and adjust the totals."

TIP You can iterate on *part* of an answer: "Keep everything, but rewrite only the closing paragraph — it fizzles out." Claude will preserve the rest. This is far faster than regenerating the whole thing and hoping.

Precision Moves Worth Memorizing

These short follow-ups fix the most common problems in one line:

- **"More concrete — add examples and numbers."** (cures abstract waffle)
- **"You're being too agreeable. What's weak about this?"** (cures flattery)
- **"Simplify: explain it like I'm smart but new to this."** (cures jargon)
- **"Give me 3 noticeably different versions, not 3 variations of one idea."** (cures samey options)
- **"Before rewriting, tell me what you think I'm asking for."** (cures talking past each other)
- **"Half the length. Same information."** (cures bloat)

Edit the Prompt vs. Reply — Know the Difference

You have two ways to iterate, and picking the right one matters:

Reply below (continue the conversation) when you're building on the answer: refining, extending, asking follow-ups. The history stays and Claude remembers it.

Edit your original message (hover over your message and edit it, then resend) when the prompt itself was flawed — you forgot key context or asked the wrong question. Editing creates a fresh branch, as if the bad exchange never happened, which keeps the conversation clean instead of dragging a wrong turn along behind you.

REMEMBER If you find yourself typing "no, ignore that, actually I meant..." more than once in a row, stop replying and edit the original prompt instead. Correcting a misunderstanding three replies deep is like renovating on a cracked foundation.

WARNING Iteration has diminishing returns. If you've looped 4-5 times and it's still not right, the problem is almost never Claude's stubbornness — it's that the prompt is missing information or an example. Go back, add the missing block from Chapter 14 (usually context or an example), and start clean.

CHAPTER 16

Power Techniques: The Pro Moves

The six blocks from Chapter 14 will carry you far. These techniques are the next gear — each one is a small habit that produces outsized results.

Technique 1: Make Claude Interview You First

The single most underused move in all of prompting. When a task is complex or you're not sure what to include, flip the script:

I need to write a business plan for a food truck. Before you write anything, ask me the 5-8 questions you need answered to do this well. Then wait for my answers.

Claude asks sharp questions (budget? city? cuisine? experience?), you answer in quick fragments, and the final output is built on exactly the right information. This turns Claude from a guesser into a consultant.

TIP Works brilliantly for anything personal: gift ideas, trip plans, career advice, workout routines. You'll be surprised what Claude asks that you'd never have thought to mention.

Technique 2: Ask for Reasoning Before Conclusions

For anything analytical — decisions, math-y questions, tricky tradeoffs — ask Claude to reason first:

Think through this step by step before giving your recommendation: should I lease or buy a car, given that I drive 25,000 km/year and keep cars about 8 years?

Asking for step-by-step thinking makes the logic visible so you can *check* it — and catch the step where an assumption doesn't match your reality.

TECHNICAL STUFF This is called *chain-of-thought prompting*. Modern Claude models also have extended-thinking modes that do deep reasoning internally, but explicitly asking for visible steps remains valuable precisely because you can audit the reasoning, not just the conclusion.

Technique 3: Label Your Prompt's Parts

When a prompt mixes instructions with pasted material, Claude can occasionally blur where one ends and the other begins. Fence things off with simple labels:

Summarize the customer email below in 2 sentences, then draft a reply that offers a refund.

EMAIL:

"""

(pasted email here)

"""

REPLY STYLE EXAMPLE:

"""

(a past reply you liked)

"""

Triple quotes, dashes, or labeled sections all work. Anthropic's docs suggest XML-style tags like `<email>...</email>` for the same purpose — Claude is specifically trained to respect those fences.

Technique 4: Demand Difference, Not Variations

Ask for "5 ideas" and you often get one idea wearing five hats. Force genuine spread:

Give me 5 approaches that are genuinely different from each other — different strategy, not different wording. For each, one sentence on who it's best for.

Or steer the spread yourself: "one safe option, one bold option, one weird option I probably haven't considered."

Technique 5: Set the Knowledge Dial

Tell Claude what you already know so it doesn't waste your time re-explaining or, worse, talking over your head:

- "I'm a complete beginner — assume nothing."
- "I know Excel well but I've never used pivot tables."
- "Skip the basics; I want the advanced considerations."

One sentence, and every explanation after it lands at the right altitude.

Technique 6: Ask for Confidence and Sources

For factual matters, add a calibration request:

Where you're confident, say so plainly. Where you're uncertain or things may have changed recently, flag it. If this depends on my country, ask which country first.

You can also simply say "search the web for the current numbers" when recency matters — Claude with web search enabled will check instead of relying on training data.

Technique 7: The Pre-Mortem

Before acting on any plan Claude helped you make, run this:

```
Assume this plan failed six months from now. What are the
3 most likely reasons why? What would I wish I'd done
differently?
```

Ten seconds of typing; routinely surfaces the blind spot you were about to walk into.

Technique 8: Let Claude Write Your Prompts

The ultimate meta-move. If a prompt matters — one you'll reuse weekly, or one for a high-stakes task — draft it *with* Claude:

```
I want to create a reusable prompt that turns my messy
meeting notes into clean minutes with action items. Write
that prompt for me, then show me how it performs on the
sample notes below.
```

Claude knows what Claude responds well to. Prompts written this way tend to be better structured than what most of us write by hand — and you can save the result as a template (Chapter 19).


CHAPTER 17

Conversation Hygiene: Keeping Claude Sharp

The Context Window Is a Whiteboard, Not a Brain

From Part I you know Claude's "working memory" — the context window — holds your whole current conversation. Here's the efficiency angle: **how you manage that whiteboard directly affects answer quality and speed.**

Figure 17-1: Why marathon chats go stale



Everything in the chat *influences* Claude — including things you've mentally moved past. That draft you rejected an hour ago? Still on the whiteboard, still gently pulling answers toward it.

The One-Mission Rule

TIP **One chat = one mission.** Planning a trip? One chat. Debugging a spreadsheet? New chat. Writing a toast for a wedding? New chat. Missions can be long — a whole project's worth of back-and-forth — but when the *topic* changes, the chat should too.

Starting fresh costs you five seconds. Continuing in a polluted chat costs you quality on every single answer after.

Start a new chat when: - The topic changes (obviously) - You've iterated into a dead end and want a genuinely fresh take - The chat has gotten so long that responses feel slower or drift back to abandoned ideas - You gave experimental instructions ("answer only in pirate speak") that you don't want haunting you

Stay in the same chat when: - You're building on earlier work and the history *is* the context - Claude needs to remember decisions made 20 messages ago

The Handoff Summary: Moving Between Chats

The pro move for ending a long working session — ask Claude to pack its own suitcase:

Summarize this conversation as a briefing for a new chat: the goal, key decisions made, current draft status, and what's left to do. I'll paste it into a fresh conversation.

Paste that briefing into a new chat and you've kept 100% of the useful context and 0% of the clutter. (If you use paid features like past-chats search or Projects, Claude can also reference earlier conversations — but the handoff summary works everywhere, on every plan, and forces useful clarity about where the work actually stands.)

REMEMBER Long documents count too. If you pasted a 100-page contract to answer one question, don't keep chatting about unrelated things in that window — every response is still hauling that contract around.

WARNING In very long chats you may eventually hit the context limit — Claude starts losing the earliest messages or the app warns you the conversation is full. Don't wait for that cliff. If a mission is big, use the handoff summary *before* you're forced to.

CHAPTER 18

Projects: Your Efficiency Engine


The Problem Projects Solve

By now you've noticed the pattern behind every technique in this book: **context is everything — and retyping context is the tax you pay in every new chat.** Projects (a paid-plan feature on claude.ai) abolish that tax.

A Project is a folder with a memory. It has two superpowers:

1. **Project instructions** — standing orders that apply to every chat inside the Project.
2. **Project knowledge** — files and documents that every chat inside can see.

Figure 18-1: How a Project changes the game



Setting Up a Project That Pulls Its Weight

Step 1 — Create it around a recurring context, not a one-off task. Good Project ideas: *My Job*, *My Novel*, *Family Finances*, *Kids' School Stuff*, *The Kitchen Renovation*, *Learning Spanish*.

Step 2 — Write instructions like an onboarding doc. Cover the who/what/how that you keep repeating:

I'm Maria, I run a 3-person pottery studio in Frankfurt.
 You help me with marketing and customer emails.
 Tone: warm, playful, never salesy. German formality: "du".
 Always keep emails under 150 words.
 Our bestsellers: beginner wheel classes and date-night
 workshops. Never discount the date-night workshops.

Step 3 — Load the knowledge files Claude should always have on hand: price lists, style guides, the resume you're always tweaking, the novel draft, the syllabus. Now "check this against our price list" just *works* — no re-uploading.

Step 4 — Refine over time. Whenever you catch yourself correcting Claude the same way twice inside a Project ("I told you, no exclamation marks"), promote that correction into the Project instructions. That's how the Project gets smarter every week.

TIP Claude can write its own onboarding. Inside a chat where things are going well, ask: *"Based on everything in this conversation, draft Project instructions so future chats start with this understanding."* Paste the result into the Project settings. Done.

TECHNICAL STUFF Project knowledge occupies context, but claude.ai manages large knowledge bases intelligently (retrieving relevant parts when the knowledge exceeds what fits). Practical upshot: prefer a few well-chosen, current documents over dumping your entire hard drive in — relevance beats volume.

WARNING Stale knowledge is worse than no knowledge. If the price list changed, update the file in the Project — otherwise Claude will confidently keep quoting last year's prices in every chat. Put a quarterly "prune the Projects" reminder in your calendar.

Projects vs. Memory vs. Styles — Which Does What?

These three features overlap in beginners' minds. Here's the sorting:

Feature	What it remembers	Best for
Project instructions & knowledge	Standing orders + files, scoped to one area of your life	Recurring work with stable context
Memory (if enabled in settings)	Facts Claude picks up across your chats	Personal preferences that apply broadly
Styles	How Claude writes (tone, formality, length)	Making every answer sound the way you like
User preferences (settings)	Global instructions for all chats	Universal rules: "always use metric," "be concise"

Rule of thumb: **global habits** → **preferences/styles**; **area-specific context** → **a Project**.

CHAPTER 19

Stop Repeating Yourself: The Template Library

Why Templates Beat Talent

A template is a power prompt you write once, keep somewhere handy (a notes app, or as a Project instruction), and reuse with tiny edits. The pros aren't better typists — they've just stopped writing prompts from scratch.

Each template below is copy-paste ready. Words in [BRACKETS] are the blanks you fill.

The Everyday Ten

1. The Email Fixer

Rewrite this email. Goal: [get a refund / say no politely / ask for a deadline extension]. Tone: [firm but friendly]. Keep it under [120] words. Keep my key facts, cut everything else. Give me only the final email.

EMAIL DRAFT:
[paste]

2. The Decision Helper

Help me decide: [option A] vs [option B].
My situation: [2-3 sentences].
What matters most to me: [cost / time / low stress], in that order. Give me a comparison table, then your recommendation with the single strongest reason — and the strongest reason I might regret it.

3. The Document Digester

I've attached [a contract / report / manual]. I am [a tenant / a parent / a small-business owner], and I need to know: [what am I committing to / what do I need to do and by when]. Summarize in plain language, flag anything unusual or risky, then list the questions I should ask before signing/acting.

4. The Explainer

Explain [topic] to me. My level: [total beginner / know the basics]. Use an analogy from [cooking / sports / everyday life]. End with the 3 things worth remembering, and one common misconception people have about it.

5. The Meeting-Notes Cleaner

Turn my messy notes below into: 1) a 3-sentence summary, 2) decisions made, 3) action items as a table with columns Owner / Task / Deadline. Anything unclear, list under "Needs clarification" instead of guessing.

NOTES:
[paste]

6. The Plan Builder

Build me a plan to [goal] by [date]. Constraints: [budget / hours per week / equipment]. Before writing the plan, ask me up to 5 questions you need answered. Then give me a week-by-week plan with a realistic "minimum viable" version for weeks when life happens.

7. The Feedback Coach

Act as a tough but fair [editor / hiring manager / customer]. Here's my [draft / resume / pitch]. First tell me what works. Then the 3 highest-impact improvements, most important first, each with a concrete example of the fix. Don't rewrite the whole thing yet.

DRAFT:
[paste]

8. The Options Generator

Give me 7 ideas for [problem]. Make them genuinely different strategies, not variations. Format: one line each – the idea, then who it suits. Mark your top pick with ★ and say why in one sentence.

9. The Learning Companion

I'm learning [topic]. Quiz me with [5] questions, one at a time – wait for my answer before showing the next. After each answer, tell me if I'm right and add one insight I didn't mention. Keep score. Start now.

10. The Devil's Advocate

I'm about to [decision]. Argue AGAINST it as persuasively as you can – steelman the opposition, don't be gentle. Then tell me which of those counterarguments you find genuinely strong versus merely rhetorical.

TIP Store your templates where friction is lowest: a pinned note, a text-expander app, or — best — as Project instructions so they're simply *always in force*. The best template is one you don't even have to paste.

REMEMBER Templates are living documents. Every time a template's output disappoints, don't just fix the answer — fix the template. Six weeks of this and your library will be tuned to your life with uncanny precision.

CHAPTER 20

Workflows That Multiply Your Output

A workflow is prompts arranged in a sequence, where each step feeds the next. This is where efficiency stops being about single answers and starts being about *finished projects*.


Workflow 1: Draft → Critique → Revise (the Quality Loop)

The most useful workflow in existence, because it turns Claude into both writer *and* editor:

- Step 1: "Write a first draft of [thing]. Don't polish – prioritize getting all the ideas down."
- Step 2: "Now switch roles: you're a demanding editor. Critique that draft – what's weak, missing, or confusing? Don't rewrite yet."
- Step 3: "Revise the draft using your own critique. Highlight what changed."

Splitting *writing* from *judging* into separate steps produces dramatically better work than "write something great" ever does — the same reason human writers don't edit while drafting.

Figure 20-1: The quality loop



One prompt = one hat. Writer hat, then editor hat, then reviser hat. Never all three hats at once.

Workflow 2: The Funnel (Broad → Narrow → Deep)

For anything open-ended — naming, planning, researching, choosing:

- Step 1 (BROAD): "Give me 15 quick options for [X]."
- Step 2 (NARROW): "I like 3, 7, and 12. Develop each in a paragraph – strengths, weaknesses."
- Step 3 (DEEP): "Let's go with 7. Build the full version."

WARNING Beginners invert the funnel: they ask for one fully developed answer immediately, don't love it, and iterate expensively on a option they'd never have chosen from a lineup. Cheap breadth first, expensive depth last.

Workflow 3: The Assembly Line (for Big Deliverables)

Big deliverables — a report, a presentation, a wedding speech, a business plan — come out mediocre when requested in one shot. Break them down:

Step 1: "We're writing [deliverable] for [audience].
Propose an outline. Don't write content yet."
Step 2: You edit the outline. (This is YOUR highest-leverage minute in the whole workflow — steering is cheap at the outline stage, expensive at the draft stage.)
Step 3: "Outline approved with my changes. Write section 1."
Step 4: Review, adjust, "Now section 2." ...and so on.
Step 5: "Now read the whole thing start to finish and fix inconsistencies, repetition, and tone drift."

TIP In step 3 onward, going section by section keeps each piece within your attention span — you'll actually *read* three paragraphs and catch problems, whereas nobody carefully reads a ten-page wall delivered at once.

Workflow 4: The Second Opinion

For decisions that matter, deliberately collect disagreement:

Chat 1: "Recommend a solution for [problem]." → note answer.
Chat 2 (fresh chat, no shared history): "Here's a proposed solution to [problem]: [paste]. Find the flaws."

Because chat 2 has no memory of producing the answer, it critiques honestly instead of defending its own work. Poor man's expert panel; costs two minutes.

Workflow 5: The Recurring Job

For anything you do weekly — status reports, newsletters, meal plans, lesson prep — combine everything from Part II:

1. Build the prompt with Claude's help (Technique 8, Chapter 16).
2. Test it, refine it twice.
3. Save it as a template (Chapter 19) or Project instruction (Chapter 18).
4. Each week, paste only the new inputs: "Here are this week's numbers → produce the usual."

First week: 20 minutes of setup. Every week after: 2 minutes. That's the entire secret of the people who seem suspiciously productive with Claude.

CHAPTER 21

Working Efficiently with Files and Artifacts

Files: Ask Before You Read

The inefficient habit: read the 40-page document yourself, *then* ask Claude questions. The efficient habit: upload first, let Claude read it, and interrogate:

```
"What are the 5 things in this document that most affect me  
as [a renter / the project lead / a beginner]?"  
"Find every date and deadline and list them chronologically."  
"Compare these two versions – what actually changed?"  
"Is there anything in here that contradicts itself?"
```

TIP For spreadsheets, don't ask generic questions ("analyze this"). Name the decision: "Which product line should I stop stocking, based on this sales data? Show your reasoning." Decision-first questions get analysis; generic questions get description.

Artifacts: Iterate on the Thing Itself

When Claude builds something substantial — a document, a webpage, a small app, a chart — it appears as an Artifact in a side panel. The efficiency win: **you iterate on the artifact directly** ("make the header blue," "add a reset button," "swap sections 2 and 3") and Claude updates it in place. You're editing a living object, not re-requesting from scratch.

Great beginner uses that feel like magic: - "Make me an interactive quiz to help my kid practice multiplication tables." - "Turn this data into a chart I can play with." - "Build a simple tip-splitting calculator for dinners with friends." - "Create a one-page checklist I can print for our camping trip."

File Creation: Ask for the Format You Actually Need

Claude can produce real downloadable files — Word documents, PowerPoint decks, Excel spreadsheets, PDFs. The efficient move is to say the destination up front:

```
"...and give it to me as a Word document I can send to my  
landlord."  
"...as a PowerPoint, max 8 slides, one idea per slide."  
"...as an Excel budget with formulas, so totals update  
when I change numbers."
```

WARNING Saying the format at the *end* of a long back-and-forth ("oh, can that be a spreadsheet?") works, but saying it up front lets Claude structure the work for that format from the start — presentations get built in slide-sized ideas, spreadsheets get built with formulas in mind.

CHAPTER 22

The Seven Deadly Sins of Prompting

Every hour wasted with Claude traces back to one of these seven. Check your last frustrating chat against this list — the culprit is in here.

Sin 1: The Mind-Reader Assumption. "Write something for the thing tomorrow." Claude knows nothing you don't say. **Penance:** the 30-second brief (Chapter 13).

Sin 2: The Kitchen-Sink Chat. One chat containing taxes, a birthday poem, and a plumbing question. Old context contaminates new answers. **Penance:** one chat, one mission (Chapter 17).

Sin 3: Accepting the First Draft. Taking answer #1 as final, when answer #2 — one line of feedback away — is where the good stuff lives. **Penance:** the iteration loop (Chapter 15). Budget for one loop, always.

Sin 4: The Vague Complaint. "Meh, try again." Claude changes random things and you spin in circles. **Penance:** diagnose + direct (Chapter 15).

Sin 5: Never Showing an Example. Three paragraphs describing the tone you want, when pasting one sample would nail it. **Penance:** Block 6 (Chapter 14).

Sin 6: Blind Trust. Shipping facts, figures, legal claims, or code without checking. Claude is superb and *still* confidently wrong sometimes — that's not a contradiction, it's the deal. **Penance:** for anything with stakes, verify names, numbers, dates, and claims — or ask Claude to search the web for current sources and check them yourself.

Sin 7: Under-Delegation. Asking for crumbs ("give me an opening line") when Claude could take the whole loaf ("here are my notes — draft the entire thing, then tell me what you need from me"). **Penance:** delegate outcomes (Chapter 13).

REMEMBER Sins 1-5 and 7 waste your *time*. Sin 6 is the only one that can genuinely *hurt* you. If you internalize a single warning from this entire book: **Claude rows, you steer, and you double-check anything that matters before it leaves your hands.**

CHAPTER 23

The One-Page Cheat Sheet

Print this. Pin it. In a month you won't need it.

GETTING REAL WORK OUT OF CLAUDE – CHEAT SHEET
BEFORE YOU TYPE • Brief it like a freelancer, not a search engine • Say the outcome: "...so that ____" • New topic? New chat.
THE SIX BLOCKS (use what the task needs) ROLE → CONTEXT → TASK → FORMAT → CONSTRAINTS → EXAMPLE Minimum viable prompt = context + clear task
MAGIC PHRASES "Ask me 5 questions before you start." "Think step by step, then recommend." "Give me only the final text, no explanation." "Make the options genuinely different." "Act as my critic – what's weak here?" "Half the length, same information."
WHEN THE ANSWER ISN'T RIGHT 1. Say what's wrong + what you want instead 2. Misunderstood from the start? EDIT the prompt 3. Stuck after 4 tries? Missing context or example
BIG TASKS Outline → approve → section by section → final pass Draft → critique → revise (separate steps!) 15 quick options → develop 3 → build 1
STOP REPEATING YOURSELF Same correction twice? → template / Project instruction Same context weekly? → make it a Project
THE ONE RULE THAT ISN'T OPTIONAL Verify anything with stakes: names, numbers, money, law, health, code. Claude rows. YOU steer.

Coming in Volume 2

You now know how to *use* Claude efficiently. Volume 2 is about *building* with it: Claude Code for turning ideas into real software, the API for automating workflows, connectors and MCP for wiring Claude into your email, calendar, and tools, agents that work while you sleep — and the advanced techniques that make all of it reliable.

Until then: go delegate something.

Product features, plans, and limits mentioned in this book (Projects, memory, styles, file creation, web search) evolve quickly — for current details, see the official documentation at support.claude.com and docs.claude.com.

— End of Part II —